

**Problem 1**

The R function to find the mean and standard deviation is

```
#Function for Problem 1
#Purpose is to find mean and standard deviation of all columns
#Mean is stored in column 1 of output
#Standard Deviation is stored in column 2 of output

Problem1Fun<-function(InTable){
  OutVec<-matrix(c(1:2*dim(InTable)[2]), ncol=2, nrow=dim(InTable)[2])
  for(i in 1:dim(InTable)[2]){
    OutVec[i,1]<-mean(InTable[,i])
    OutVec[i,2]<-sd(InTable[,i])
  }
  return(OutVec)
}
```

The output for running `Problem1Fun(Nitrates)` is

```
> Problem1Fun(Nitrates)
      [,1]      [,2]
[1,]  2.00000 1.351725e+00
[2,] 33.68870 8.366869e-02
[3,] -112.21931 3.152602e-01
[4,]  483.08873 1.128625e+02
[5,] 37431.11734 1.647951e+04
[6,]  679.43042 1.093556e+03
[7,] 26499.28471 1.576737e+04
[8,]  26.52196 2.216925e+01
[9,]  17.79275 7.931085e+01
[10,]  7.54878 5.698698e-01
[11,] 613.55278 1.320502e+03
```

To use the `apply` function to produce the means in 1 line, we would use the following command and the output is as follows

```
> apply(Nitrates,2,mean)
      LAND.USE.LEVEL1      latitude_decimals      longitude_decimal
      2.00000      33.68870      -112.21931
      elevation.Meters. DistanceFromUrbanCenter      POP.DENSITY
      483.08873      37431.11734      679.43042
      INCOME.PER.CAPITA      VegCover      NO3.N
      26499.28471      26.52196      17.79275
      pH      conductivity
      7.54878      613.55278
```

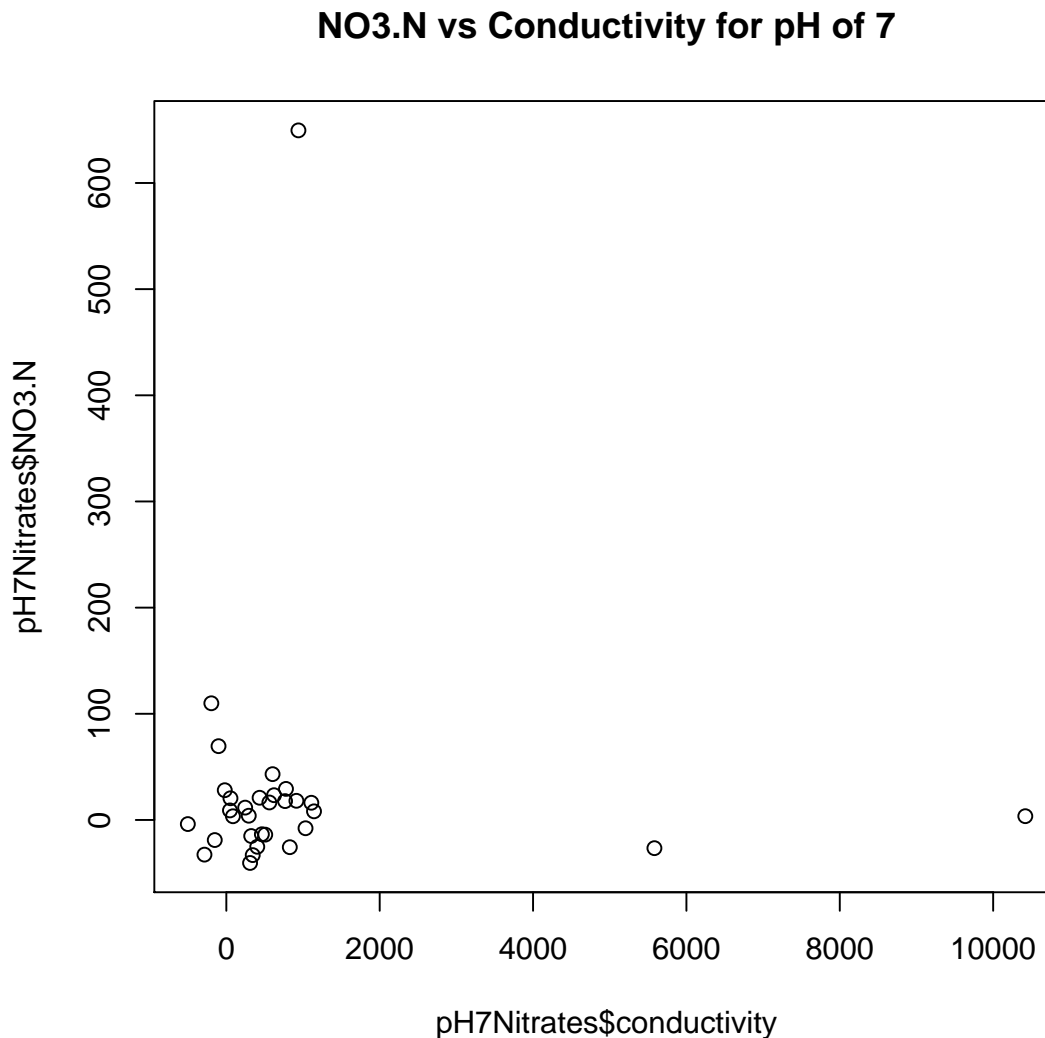
Similarly, we have the following for standard deviation:

```
> apply(Nitrates,2,sd)
      LAND.USE.LEVEL1      latitude_decimals      longitude_decimal
      1.351725e+00      8.366869e-02      3.152602e-01
      elevation.Meters. DistanceFromUrbanCenter      POP.DENSITY
      1.128625e+02      1.647951e+04      1.093556e+03
      INCOME.PER.CAPITA      VegCover      NO3.N
      1.576737e+04      2.216925e+01      7.931085e+01
      pH      conductivity
      5.698698e-01      1.320502e+03
```

**Problem 2**

We create a dataframe with rows with a pH value of 7, then plot NO3 against conductivity with the following code and then the output plot is the included



```
> pH7Nitrates<-Nitrates[which(Nitrates$pH==7),]  
> plot(pH7Nitrates$conductivity,pH7Nitrates$NO3.N,  
       main="NO3.N vs Conductivity for pH of 7")  
> dev.print(device=pdf,"Problem2.pdf")
```

**Problem 3**

To produce a side-by-side histogram of NO3.N for pH of 7 and pH of 8, we create two dataframes and then use `hist()` which produces the subsequent plot.

```
> pH7Nitrates<-Nitrates[which(Nitrates$pH==7),]  
> pH8Nitrates<-Nitrates[which(Nitrates$pH==8),]  
> par(mfrow=c(1,2))  
> hist(pH7Nitrates$NO3.N)  
> hist(pH8Nitrates$NO3.N)  
> dev.print(device=pdf,"Problem3.pdf")
```

## Histogram of pH7Nitrates\$NO3    Histogram of pH8Nitrates\$NO3



### Problem 4

To determine how many data points have a landuse of 2, we use the following line

```
> dim(Nitrates[which(Nitrates$LAND.USE.LEVEL1==2),])[1]
[1] 45
```

So, we see that there are 45 data points with a landuse of 2.

### Problem 5

After installing packages with `install.packages('mvtnorm')` then using the following given code

```
library(mvtnorm)

for(i in 1:1000){
  currentrealization<-rmvnorm(1,c(0,0),CovMat)
  Realizations[i,1]<-currentrealization[1,1]
  Realizations[i,2]<-currentrealization[1,2]
}
```

we can create a covariance matrix with the given values and create a matrix of zeros

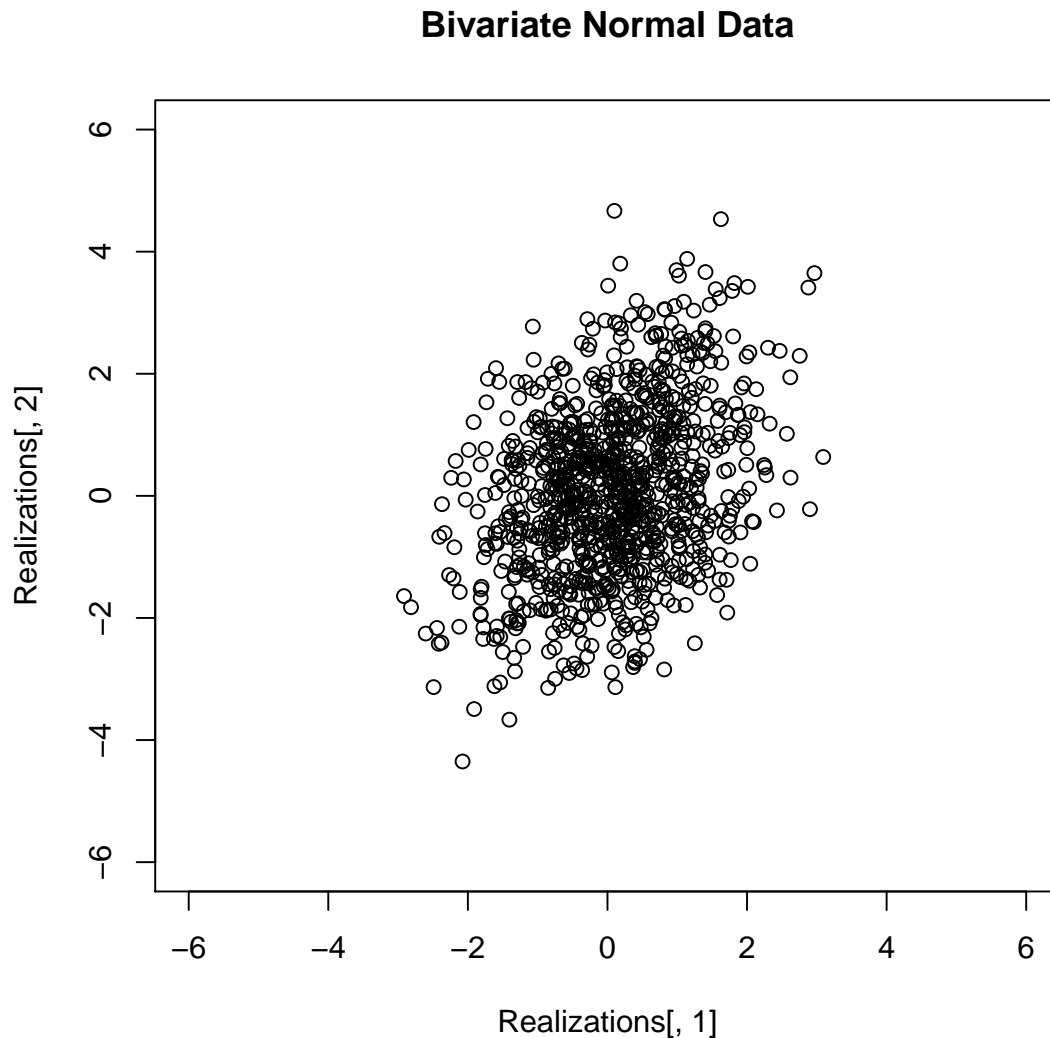
```
> CovMat<-matrix(rep(1,4),nrow=2,ncol=2)
> CovMat[1,]<-c(1,0.5)
> CovMat[2,]<-c(0.5,2)
> CovMat
      [,1] [,2]
[1,]  1.0  0.5
[2,]  0.5  2.0

> Realizations<-matrix(0,nrow=1000,ncol=2)
```

then using the given code, we have the following scatter-plot by


```
> plot(Realizations[,1],Realizations[,2],main="Bivariate Normal Data",
      xlim=c(-6,6),ylim=c(-6,6))
> dev.print(device=pdf,"Problem5.pdf")
```

For this data we have the following as the means of the columns



```
> apply(Realizations,2,mean)
[1] -0.03239635 -0.03260434
```

We then run the problem again and we have the following scatter-plot and for this we have the following as



the means of the columns

```
> apply(Realizations,2,mean)
[1] 0.01286922 0.03089901
```

As we see for the means, both are close to 0, and as we ran the code for a second time, we see that (perhaps by co-incidence) we have means closer, in magnitude, to 0.

#### Problem 6

Using `rexp` we see the syntax of inputs is `rexp(n,1/mean)`, where `n` is the number of outcomes. So, for a mean of 1 and 50 outcomes, we can use `rexp(50,1)`. Thus, to produce the sum, we would use

```
> sum(rexp(50,1))
```

To produce this sum 1000 times, we use the following loop

```
realization<-rep(1,1000)
for( i in 1:1000){
    realization[i]<-sum(rexp(50,1))
}
```

Then using `length`, we have the percentage of sums above 60

```
> 100*length(which(realization >60))/length(realization)
[1] 8.8
```

Thus there are 8.8% of the 1000 sums above 60.

### Problem 7

To find the number of iterations it takes to find the sum of 50 realizations over 80, we run

```
i<-0
testsum<80
while(testsum<80){
    testsum<-sum(rexp(50,1))
    i<-i+1
}
print(i)
```

Running this a few times, we require 1554 iterations, 16034 iterations, 3313 iterations to have the sum over 80.

### Problem 8

Using the following code

```
> X<-rnorm(1000,0,1)
> Y<-rnorm(1000,0,2)
> Z<-X+Y
> hist(Z)
> dev.print(device=pdf,"Problem8.pdf")
```

we have the following histogram This plot seems to be a Normal distribution about 0.

### Problem 9 Using the following code

```
> X<-rnorm(1000,0,1)
> Y<-rnorm(1000,X,2)
> Z<-X+Y
> hist(Z)
> dev.print(device=pdf,"Problem9.pdf")
```

we have the following plot. Which still looks Normal about 0, but shows a bit of being skewed.




Figure 1: Problem 8 Plot




Figure 2: Problem 9 Plot